

Collections Framework

The java.util package also contains one of Java's most powerful subsystems: the Collections Framework. The Collections Framework is a sophisticated hierarchy of interfaces and classes that provide state-of-the-art technology for managing groups of objects. It merits close attention by all programmers. Because java.util contains a wide array of functionality, it is quite large.

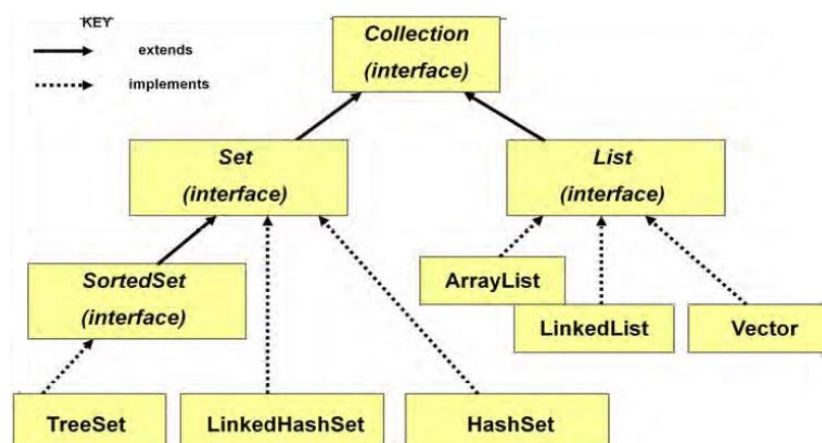
Collections Overview

The Java Collections Framework standardizes the way in which groups of objects are handled by your programs.

Benefits of the Java Collections Framework

- **Reduces programming effort:** By providing useful data structures and algorithms, the Collections Framework frees you to concentrate on the important parts of your program rather than on the low-level work required to make it work.
- **Increases program speed and quality:** This Collections Framework provides high-performance, high-quality implementations of useful data structures and algorithms.
- **Allows interoperability among unrelated APIs:** APIs will interoperate seamlessly w.r.to. collections, even though they were written independently.
- **Reduces effort to learn and to use new APIs:** Many APIs naturally take collections on input and furnish them as output. With the advent of standard collection interfaces, the effort to learn and to use new APIs is reduced.
- **Reduces effort to design new APIs:** An API that relies on collections can easily be designed using standard collection interfaces.
- **Promotes software reuse:** new data structures that conform to the standard collection interfaces are by nature reusable. The same goes for new algorithms that operate on objects that implement these interfaces.

The Collection Interfaces



The Collections Framework defines several core interfaces. This section provides an overview of each interface. Beginning with the collection interfaces is necessary because they determine the fundamental nature of the collection classes. Put differently, the concrete classes simply provide different implementations of the standard interfaces. The interfaces that underpin collections are summarized in the following table:

Interface	Description
Collection	Enables you to work with groups of objects; it is at the top of the collections hierarchy.
Deque	Extends Queue to handle a double-ended queue.
List	Extends Collection to handle sequences (lists of objects).
NavigableSet	Extends SortedSet to handle retrieval of elements based on closest-match searches.
Queue	Extends Collection to handle special types of lists in which elements are removed only from the head.
Set	Extends Collection to handle sets, which must contain unique elements.
SortedSet	Extends Set to handle sorted sets.

In addition to the collection interfaces, collections also use the **Comparator**, **RandomAccess**, **Iterator**, and **ListIterator** interfaces. Beginning with JDK 8, **Spliterator** can also be used. Briefly, **Comparator** defines how two objects are compared; **Iterator**, **ListIterator**, and **Spliterator** enumerate the objects within a collection. By implementing **RandomAccess**, a list indicates that it supports efficient, random access to its elements. To provide the greatest flexibility in their use, the collection interfaces allow some methods to be optional. The optional methods enable you to modify the contents of a collection. Collections that support these methods are called modifiable. Collections that do not allow their contents to be changed are called unmodifiable. If an attempt is made to use one of these methods on an unmodifiable collection, an **UnsupportedOperationException** is thrown. All the built-in collections are modifiable. The following sections examine the collection interfaces.

The Collection Interface

The **Collection** interface is the foundation upon which the Collections Framework is built because it must be implemented by any class that defines a collection. **Collection** is a generic interface that has this declaration:

interface Collection <E>

Here, **E** specifies the type of objects that the collection will hold. **Collection** extends the **Iterable** interface. This means that all collections can be cycled through by use of the for each style for loop. (Recall that only classes that implement **Iterable** can be cycled through by the for.)

Collection declares the core methods that all collections will have. These methods are summarized in Table 18-1. Several of these methods can throw an **UnsupportedOperationException**. This occurs if a collection cannot be modified. A **ClassCastException** is generated when one object is incompatible with another, such as when an attempt is made to add an incompatible object to a collection. A **NullPointerException** is thrown if an attempt is made to store a null object and null elements are not allowed in the collection. An **IllegalArgumentException** is thrown if an invalid argument is used. An **IllegalStateException** is thrown if an attempt is made to add an element to a fixed-length collection that is full.

Method	Description
<code>boolean add(E obj)</code>	Adds <i>obj</i> to the invoking collection. Returns true if <i>obj</i> was added to the collection. Returns false if <i>obj</i> is already a member of the collection and the collection does not allow duplicates.
<code>boolean addAll(Collection<? extends E> c)</code>	Adds all the elements of <i>c</i> to the invoking collection. Returns true if the collection changed (i.e., the elements were added). Otherwise, returns false .
<code>void clear()</code>	Removes all elements from the invoking collection.
<code>boolean contains(Object obj)</code>	Returns true if <i>obj</i> is an element of the invoking collection. Otherwise, returns false .
<code>boolean containsAll(Collection<?> c)</code>	Returns true if the invoking collection contains all elements of <i>c</i> . Otherwise, returns false .
<code>boolean equals(Object obj)</code>	Returns true if the invoking collection and <i>obj</i> are equal. Otherwise, returns false .
<code>int hashCode()</code>	Returns the hash code for the invoking collection.
<code>boolean isEmpty()</code>	Returns true if the invoking collection is empty. Otherwise, returns false .
<code>Iterator<E> iterator()</code>	Returns an iterator for the invoking collection.
<code>default Stream<E> parallelStream()</code>	Returns a stream that uses the invoking collection as its source for elements. If possible, the stream supports parallel operations. (Added by JDK 8.)
<code>boolean remove(Object obj)</code>	Removes one instance of <i>obj</i> from the invoking collection. Returns true if the element was removed. Otherwise, returns false .
<code>boolean removeAll(Collection<?> c)</code>	Removes all elements of <i>c</i> from the invoking collection. Returns true if the collection changed (i.e., elements were removed). Otherwise, returns false .
<code>default boolean removeIf(Predicate<? super E> predicate)</code>	Removes from the invoking collection those elements that satisfy the condition specified by <i>predicate</i> . (Added by JDK 8.)

Table 18-1 The Methods Declared by **Collection**

Method	Description
<code>boolean retainAll(Collection<?> c)</code>	Removes all elements from the invoking collection except those in <i>c</i> . Returns true if the collection changed (i.e., elements were removed). Otherwise, returns false .
<code>int size()</code>	Returns the number of elements held in the invoking collection.
<code>default Spliterator<E> spliterator()</code>	Returns a spliterator to the invoking collections. (Added by JDK 8.)
<code>default Stream<E> stream()</code>	Returns a stream that uses the invoking collection as its source for elements. The stream is sequential. (Added by JDK 8.)
<code>Object[] toArray()</code>	Returns an array that contains all the elements stored in the invoking collection. The array elements are copies of the collection elements.
<code><T> T[] toArray(T array[])</code>	Returns an array that contains the elements of the invoking collection. The array elements are copies of the collection elements. If the size of <i>array</i> equals the number of elements, these are returned in <i>array</i> . If the size of <i>array</i> is less than the number of elements, a new array of the necessary size is allocated and returned. If the size of <i>array</i> is greater than the number of elements, the array element following the last collection element is set to null . An ArrayStoreException is thrown if any collection element has a type that is not a subtype of <i>array</i> .

Table 18-1 The Methods Declared by **Collection** (continued)

Note: Any 6

Objects are added to a collection by calling `add()`. Notice that `add()` takes an argument of type `E`, which means that objects added to a collection must be compatible with the type of data expected by the collection.

- You can add the entire contents of one collection to another by calling `addAll()`.
- You can remove an object by using `remove()`.
- To remove a group of objects, call `removeAll()`.
- You can remove all elements except those of a specified group by calling `retainAll()`.
- Beginning with JDK 8, to remove an element only if it satisfies some condition, you can use `removeIf()`.
- To empty a collection, call `clear()`.
- You can determine whether a collection contains a specific object by calling `contains()`.
- To determine whether one collection contains all the members of another, call `containsAll()`.
- You can determine when a collection is empty by calling `isEmpty()`.
- The number of elements currently held in a collection can be determined by calling `size()`.
- The `toArray()` methods return an array that contains the elements stored in the invoking collection. The first returns an array of `Object`. The second returns an array of elements that have the same type as the array specified as a parameter. Normally, the second form is more convenient because it returns the desired array type. These methods are more important than it might at first seem. Often, processing the contents of a collection by using array-like syntax is advantageous. By providing a pathway between collections and arrays, you can have the best of both worlds.
- Two collections can be compared for equality by calling `equals()`. The precise meaning of “equality” may differ from collection to collection. For example, you can implement `equals()` so that it compares the values of elements stored in the collection. Alternatively, `equals()` can compare references to those elements.
- Another important method is `iterator()`, which returns an iterator to a collection.
- The new `splititerator()` method returns a `splititerator` to the collection.
- Iterators are frequently used when working with collections.
- Finally, the `stream()` and `parallelStream()` methods return a `Stream` that uses the collection as a source of elements

The List Interface

The `List` interface extends `Collection` and declares the behavior of a collection that stores a sequence of elements. Elements can be inserted or accessed by their position in the list, using a zero-based index. A list may contain duplicate elements. `List` is a generic interface that has this declaration:

interface List<E>

Here, `E` specifies the type of objects that the list will hold.

In addition to the methods defined by `Collection`, `List` defines some of its own, which are summarized in Table 18-2. Note again that several of these methods will throw an **`UnsupportedOperationException`** if the list cannot be modified, and a **`ClassCastException`** is generated when one object is incompatible with another, such as when an attempt is made to add an incompatible object to a list. Also, several methods will throw an **`IndexOutOfBoundsException`** if an invalid index is used. A **`NullPointerException`** is thrown if an attempt is made to store a null object and null elements are not allowed in the list. An **`IllegalArgumentException`** is thrown if an invalid argument is used.

- To the versions of `add()` and `addAll()` defined by `Collection`, `List` adds the methods `add(int, E)` and `addAll(int, Collection)`. These methods insert elements at the specified index. Also, the

semantics of `add(E)` and `addAll(Collection)` defined by `Collection` are changed by `List` so that they add elements to the end of the list.

- You can modify each element in the collection by using `replaceAll()`.
- To obtain the object stored at a specific location, call `get()` with the index of the object.
- To assign a value to an element in the list, call `set()`, specifying the index of the object to be changed. To find the index of an object, use `indexOf()` or `lastIndexOf()`.
- You can obtain a sublist of a list by calling `subList()`, specifying the beginning and ending indexes of the sublist. As you can imagine, `subList()` makes list processing quite convenient.
- One way to sort a list is with the `sort()` method defined by `List`.

Method	Description
<code>void add(int index, E obj)</code>	Inserts <i>obj</i> into the invoking list at the index passed in <i>index</i> . Any preexisting elements at or beyond the point of insertion are shifted up. Thus, no elements are overwritten.
<code>boolean addAll(int index, Collection<? extends E> c)</code>	Inserts all elements of <i>c</i> into the invoking list at the index passed in <i>index</i> . Any preexisting elements at or beyond the point of insertion are shifted up. Thus, no elements are overwritten. Returns true if the invoking list changes and returns false otherwise.
<code>E get(int index)</code>	Returns the object stored at the specified index within the invoking collection .
<code>int indexOf(Object obj)</code>	Returns the index of the first instance of <i>obj</i> in the invoking list. If <i>obj</i> is not an element of the list, -1 is returned.
<code>int lastIndexOf(Object obj)</code>	Returns the index of the last instance of <i>obj</i> in the invoking list. If <i>obj</i> is not an element of the list, -1 is returned.
<code>ListIterator<E> listIterator()</code>	Returns an iterator to the start of the invoking list.
<code>ListIterator<E> listIterator(int index)</code>	Returns an iterator to the invoking list that begins at the specified <i>index</i> .
<code>E remove(int index)</code>	Removes the element at position <i>index</i> from the invoking list and returns the deleted element. The resulting list is compacted. That is, the indexes of subsequent elements are decremented by one.
default void <code>replaceAll(UnaryOperator<E> opToApply)</code>	Updates each element in the list with the value obtained from the <i>opToApply</i> function. (Added by JDK 8.)
<code>E set(int index, E obj)</code>	Assigns <i>obj</i> to the location specified by <i>index</i> within the invoking list. Returns the old value.
default void <code>sort(Comparator<? super E> comp)</code>	Sorts the list using the comparator specified by <i>comp</i> . (Added by JDK 8.)
<code>List<E> subList(int start, int end)</code>	Returns a list that includes elements from <i>start</i> to <i>end</i> -1 in the invoking list. Elements in the returned list are also referenced by the invoking object.

Table 18-2 The Methods Declared by **List**

Note: Any 5

The Set Interface

The Set interface defines a set. It extends Collection and specifies the behavior of a collection that does not allow duplicate elements. Therefore, the add() method returns false if an attempt is made to add duplicate elements to a set. It does not specify any additional methods of its own. Set is a generic interface that has this declaration:

interface Set<E>

Here, E specifies the type of objects that the set will hold.

The Collection Classes


Now that you are familiar with the collection interfaces, you are ready to examine the standard classes that implement them. Some of the classes provide full implementations that can be used as-is. Others are abstract, providing skeletal implementations that are used as starting points for creating concrete collections. As a general rule, the collection classes are not synchronized, but as you will see later in this chapter, it is possible to obtain synchronized versions

Class	Description
AbstractCollection	Implements most of the Collection interface.
AbstractList	Extends AbstractCollection and implements most of the List interface.
AbstractQueue	Extends AbstractCollection and implements parts of the Queue interface.
AbstractSequentialList	Extends AbstractList for use by a collection that uses sequential rather than random access of its elements.
LinkedList	Implements a linked list by extending AbstractSequentialList .
ArrayList	Implements a dynamic array by extending AbstractList .
ArrayDeque	Implements a dynamic double-ended queue by extending AbstractCollection and implementing the Deque interface.
AbstractSet	Extends AbstractCollection and implements most of the Set interface.
EnumSet	Extends AbstractSet for use with enum elements.
HashSet	Extends AbstractSet for use with a hash table.
LinkedHashSet	Extends HashSet to allow insertion-order iterations.
PriorityQueue	Extends AbstractQueue to support a priority-based queue.
TreeSet	Implements a set stored in a tree. Extends AbstractSet .

The following sections examine the concrete collection classes and illustrate their use.

NOTE In addition to the collection classes, several legacy classes, such as Vector, Stack, and Hashtable, have been reengineered to support collections.

The ArrayList Class



The ArrayList class extends AbstractList and implements the List interface. ArrayList is a generic class that has this declaration:

class ArrayList<E>

Here, E specifies the type of objects that the list will hold.

ArrayList supports dynamic arrays that can grow as needed. In Java, standard arrays are of a fixed length. After arrays are created, they cannot grow or shrink, which means that you must know in advance how many elements an array will hold. But, sometimes, you may not know until run time precisely how large an array you need. To handle this situation, the Collections Framework defines ArrayList.

In essence, an ArrayList is a variable-length array of object references. That is, an ArrayList can dynamically increase or decrease in size.

Array lists are created with an initial size. When this size is exceeded, the collection is automatically enlarged. When objects are removed, the array can be shrunk.

ArrayList has the constructors shown here:

ArrayList()

ArrayList(Collection<? extends E> c)

ArrayList(int capacity)

The first constructor builds an empty array list. The second constructor builds an array list that is initialized with the elements of the collection c. The third constructor builds an array list that has the specified initial capacity. The capacity is the size of the underlying array that is used to store the elements. The capacity grows automatically as elements are added to an array list.

The following program shows a simple use of ArrayList. An array list is created for objects of type String, and then several strings are added to it. (Recall that a quoted string is translated into a String object.) The list is then displayed. Some of the elements are removed and the list is displayed again.

```
// Demonstrate ArrayList.
import java.util.*;
class ArrayListDemo {
    public static void main(String args[]) {
```

```

// Create an array list.
ArrayList<String> al = new ArrayList<String>();
System.out.println("Initial size of al: " + al.size());
// Add elements to the array list.
al.add("C");
al.add("A");
al.add("E");
al.add("B");
al.add("D");
al.add("F");
al.add(1, "A2");
System.out.println("Size of al after additions: " + al.size());
// Display the array list.
System.out.println("Contents of al: " + al);
// Remove elements from the array list.
al.remove("F");
al.remove(2);
System.out.println("Size of al after deletions: " + al.size());
System.out.println("Contents of al: " + al);
}
}

```

The output from this program is shown here:

```

Initial size of al: 0
Size of al after additions: 7
Contents of al: [C, A2, A, E, B, D, F]
Size of al after deletions: 5
Contents of al: [C, A2, E, B, D]

```

Obtaining an Array from an ArrayList

When working with ArrayList, you will sometimes want to obtain an actual array that contains the contents of the list. You can do this by calling `toArray()`, which is defined by Collection.

Several reasons exist why you might want to convert a collection into an array, such as:

- To obtain faster processing times for certain operations
- To pass an array to a method that is not overloaded to accept a collection
- To integrate collection-based code with legacy code that does not understand collections

```

Object[] toArray()
<T> T[] toArray(T array[])

```

The first returns an array of Object. The second returns an array of elements that have the same type as T. Normally, the second form is more convenient because it returns the proper type of array. The following program demonstrates its use:

```

// Convert an ArrayList into an array.
import java.util.*;
class ArrayListToArray {
    public static void main(String args[]) {
        // Create an array list.
        ArrayList<Integer> al = new ArrayList<Integer>();
        // Add elements to the array list.

```



```

al.add(1);
al.add(2);
al.add(3);
al.add(4);
System.out.println("Contents of al: " + al);
// Get the array.
Integer ia[] = new Integer[al.size()];
ia = al.toArray(ia);
int sum = 0;
// Sum the array.
for(int i : ia) sum += i;
System.out.println("Sum is: " + sum);
}
}

```

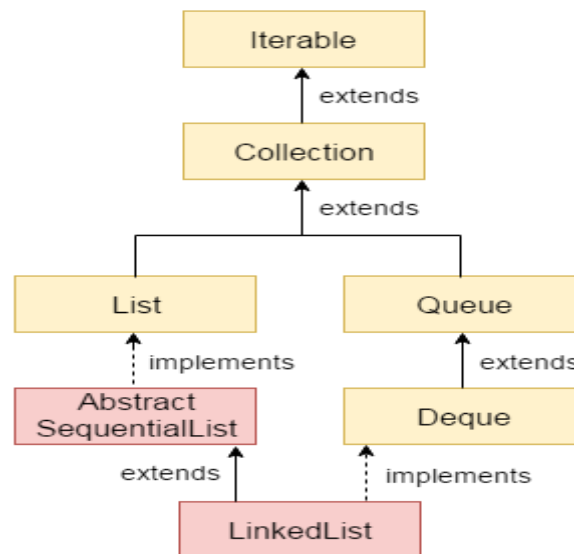
The output from the program is shown here:

Contents of al: [1, 2, 3, 4]

Sum is: 10

The program begins by creating a collection of integers. Next, `toArray()` is called and it obtains an array of `Integer`s. Then, the contents of that array are summed by use of a for-each style for loop.

The LinkedList Class



The `LinkedList` class extends `AbstractSequentialList` and implements the `List`, `Deque`, and `Queue` interfaces. It provides a linked-list data structure. `LinkedList` is a generic class that has this declaration:

```
class LinkedList<E>
```

Here, `E` specifies the type of objects that the list will hold. `LinkedList` has the two constructors shown here:

```
LinkedList()
```

```
LinkedList(Collection<? extends E> c)
```

The first constructor builds an empty linked list. The second constructor builds a linked list that is initialized with the elements of the collection c.

Because LinkedList implements the Deque interface, you have access to the methods defined by Deque. For example, to add elements to the start of a list, you can use **addFirst()** or **offerFirst()**. To add elements to the end of the list, use **addLast()** or **offerLast()**. To obtain the first element, you can use **getFirst()** or **peekFirst()**. To obtain the last element, use **getLast()** or **peekLast()**. To remove the first element, use **removeFirst()** or **pollFirst()**. To remove the last element, use **removeLast()** or **pollLast()**.

The following program illustrates LinkedList:

```
// Demonstrate LinkedList.
import java.util.*;
class LinkedListDemo {
    public static void main(String args[]) {
        // Create a linked list.
        LinkedList<String> ll = new LinkedList<String>();
        // Add elements to the linked list.
        ll.add("F");
        ll.add("B");
        ll.add("D");
        ll.add("E");
        ll.add("C");
        ll.addLast("Z");
        ll.addFirst("A");
        ll.add(1, "A2");
        System.out.println("Original contents of ll: " + ll);
        // Remove elements from the linked list.
        ll.remove("F");
        ll.remove(2);
        System.out.println("Contents of ll after deletion: " + ll);
        // Remove first and last elements.
        ll.removeFirst();
        ll.removeLast();
        System.out.println("ll after deleting first and last: " + ll);
        // Get and set a value.
        String val = ll.get(2);
        ll.set(2, val + " Changed");
        System.out.println("ll after change: " + ll);
    }
}
```

The output from this program is shown here:

Original contents of ll: [A, A2, F, B, D, E, C, Z]

Contents of ll after deletion: [A, A2, D, E, C, Z]

ll after deleting first and last: [A2, D, E, C]

ll after change: [A2, D, E Changed, C]

The HashSet Class

HashSet extends AbstractSet and implements the Set interface. It creates a collection that uses a hash table for storage. HashSet is a generic class that has this declaration:

class HashSet<E>

Here, E specifies the type of objects that the set will hold.

A hash table stores information by using a mechanism called hashing. In hashing, the informational content of a key is used to determine a unique value, called its hash code. The hash code is then used as the index at which the data associated with the key is stored. The transformation of the key into its hash code is performed automatically—you never see the hash code itself. The advantage of hashing is that it allows the execution time of `add()`, `contains()`, `remove()`, and `size()` to remain constant even for large sets.

The following constructors are defined:

HashSet()

HashSet(Collection<? extends E> c)

HashSet(int capacity)

HashSet(int capacity, float fillRatio)

The first form constructs a default hash set. The second form initializes the hash set by using the elements of `c`. The third form initializes the capacity of the hash set to `capacity`. (The default capacity is 16.) The fourth form initializes both the capacity and the fill ratio (also called load capacity) of the hash set from its arguments. The fill ratio must be between 0.0 and 1.0, and it determines how full the hash set can be before it is resized upward.

Specifically, when the number of elements is greater than the capacity of the hash set multiplied by its fill ratio, the hash set is expanded. For constructors that do not take a fill ratio, 0.75 is used.

HashSet does not define any additional methods beyond those provided by its super classes and interfaces. It is important to note that **HashSet does not guarantee the order of its elements**, because the process of hashing doesn't usually lend itself to the creation of sorted sets.

If you need sorted storage, then another collection, such as `TreeSet`, is a better choice. Here is an example that demonstrates `HashSet`:

```
// Demonstrate HashSet.
import java.util.*;
class HashSetDemo {
    public static void main(String args[]) {
        // Create a hash set.
        HashSet<String> hs = new HashSet<String>();
        // Add elements to the hash set.
        hs.add("Beta");
        hs.add("Alpha");
        hs.add("Eta");
        hs.add("Gamma");
        hs.add("Epsilon");
        hs.add("Omega");
        System.out.println(hs);
    }
}
```

The following is the output from this program:

[Gamma, Eta, Alpha, Epsilon, Omega, Beta]

As explained, the elements are not stored in sorted order, and the precise output may vary.

Accessing a Collection via an Iterator

Often, you will want to cycle through the elements in a collection. For example, you might want to display each element. One way to do this is to employ an iterator, which is an object that implements either the `Iterator` or the `ListIterator` interface. `Iterator` enables you to cycle through a collection, obtaining or removing elements. `ListIterator` extends `Iterator` to allow bidirectional traversal of a list, and the modification of elements. `Iterator` and `ListIterator` are generic interfaces which are declared as shown here:

```
interface Iterator<E>
```

```
interface ListIterator<E>
```

Here, `E` specifies the type of objects being iterated. The `Iterator` interface declares the methods shown in Table 18-8. The methods declared by `ListIterator` (along with those inherited from `Iterator`) are shown in Table 18-9. In both cases, operations that modify the underlying collection are optional. For example, `remove()` will throw **`UnsupportedOperationException`** when used with a read-only collection. Various other exceptions are possible.

Method	Description
default void forEachRemaining(Consumer<? super E> action)	The action specified by <i>action</i> is executed on each unprocessed element in the collection . (Added by JDK 8.)
boolean hasNext()	Returns true if there are more elements. Otherwise, returns false .
E next()	Returns the next element. Throws NoSuchElementException if there is not a next element.
default void remove()	Removes the current element. Throws IllegalStateException if an attempt is made to call remove() that is not preceded by a call to next() . The default version throws an UnsupportedOperationException .

Table 18-8 The Methods Declared by `Iterator`

Method	Description
<code>void add(E obj)</code>	Inserts <i>obj</i> into the list in front of the element that will be returned by the next call to next() .
default void <code>forEachRemaining(Consumer<? super E> action)</code>	The action specified by <i>action</i> is executed on each unprocessed element in the collection . (Added by JDK 8.)
<code>boolean hasNext()</code>	Returns true if there is a next element. Otherwise, returns false .
<code>boolean hasPrevious()</code>	Returns true if there is a previous element. Otherwise, returns false .
<code>E next()</code>	Returns the next element. A NoSuchElementException is thrown if there is not a next element.
<code>int nextIndex()</code>	Returns the index of the next element. If there is not a next element, returns the size of the list.
<code>E previous()</code>	Returns the previous element. A NoSuchElementException is thrown if there is not a previous element.
<code>int previousIndex()</code>	Returns the index of the previous element. If there is not a previous element, returns -1.
<code>void remove()</code>	Removes the current element from the list. An IllegalStateException is thrown if remove() is called before next() or previous() is invoked.
<code>void set(E obj)</code>	Assigns <i>obj</i> to the current element. This is the element last returned by a call to either next() or previous() .

Table 18-9 The Methods Provided by **ListIterator**

Using an Iterator

Before you can access a collection through an iterator, you must obtain one. Each of the collection classes provides an `iterator()` method that returns an iterator to the start of the collection. By using this iterator object, you can access each element in the collection, one element at a time. In general, to use an iterator to cycle through the contents of a collection, follow these steps:

1. Obtain an iterator to the start of the collection by calling the collection's `iterator()` method.
2. Set up a loop that makes a call to `hasNext()`. Have the loop iterate as long as `hasNext()` returns true.
3. Within the loop, obtain each element by calling `next()`.

For collections that implement **List**, you can also obtain an iterator by calling `listIterator()`. As explained, a list iterator gives you the ability to access the collection in either the forward or backward direction and lets you modify an element. Otherwise, **ListIterator** is used just like **Iterator**.

The following example implements these steps, demonstrating both the **Iterator** and **ListIterator** interfaces. It uses an **ArrayList** object, but the general principles apply to any type of collection. Of course, **ListIterator** is available only to those collections that implement the **List** interface.

```
// Demonstrate iterators.
import java.util.*;
class IteratorDemo {
    public static void main(String args[]) {
        // Create an array list.
        ArrayList<String> al = new ArrayList<String>();
        // Add elements to the array list.
```

```

al.add("C");
al.add("A");
al.add("E");
al.add("B");
al.add("D");
al.add("F");
// Use iterator to display contents of al.
System.out.print("Original contents of al: ");
Iterator<String> itr = al.iterator();
while(itr.hasNext()) {
    String element = itr.next();
    System.out.print(element + " ");
}
System.out.println();
// Modify objects being iterated.
ListIterator<String> litr = al.listIterator();
while(litr.hasNext()) {
    String element = litr.next();
    litr.set(element + "+");
}
System.out.print("Modified contents of al: ");
itr = al.iterator();
while(itr.hasNext()) {
    String element = itr.next();
    System.out.print(element + " ");
}
System.out.println();
// Now, display the list backwards.
System.out.print("Modified list backwards: ");
while(litr.hasPrevious()) {
    String element = litr.previous();
    System.out.print(element + " ");
}
System.out.println();
}
}

```

The output is shown here:

Original contents of al: C A E B D F

Modified contents of al: C+ A+ E+ B+ D+ F+

Modified list backwards: F+ D+ B+ E+ A+ C+

Pay special attention to how the list is displayed in reverse. After the list is modified, `litr` points to the end of the list. (Remember, `litr.hasNext()` returns false when the end of the list has been reached.) To traverse the list in reverse, the program continues to use `litr`, but this time it checks to see whether it has a previous element. As long as it does, that element is obtained and displayed.

The For-Each Alternative to Iterators

If you won't be modifying the contents of a collection or obtaining elements in reverse order, then the for-each version of the for loop is often a more convenient alternative to

cycling through a collection than is using an iterator. Recall that the for can cycle through any collection of objects that implement the Iterable interface. Because all of the collection classes implement this interface, they can all be operated upon by the for. The following example uses a for loop to sum the contents of a collection:

```
// Use the for-each for loop to cycle through a collection.
import java.util.*;
class ForEachDemo {
    public static void main(String args[]) {
        // Create an array list for integers.
        ArrayList<Integer> vals = new ArrayList<Integer>();
        // Add values to the array list.
        vals.add(1);
        vals.add(2);
        vals.add(3);
        vals.add(4);
        vals.add(5);
        // Use for loop to display the values.
        System.out.print("Contents of vals: ");
        for(int v : vals)
            System.out.print(v + " ");
        System.out.println();
        // Now, sum the values by using a for loop
        int sum = 0;
        for(int v : vals)
            sum += v;
        System.out.println("Sum of values: " + sum);
    }
}
```

The output from the program is shown here:

Contents of vals: 1 2 3 4 5

Sum of values: 15

As you can see, the for loop is substantially shorter and simpler to use than the iteratorbased approach. However, it can only be used to cycle through a collection in the forward direction, and you can't modify the contents of the collection.

LinkedList Vs. ArrayList

Both the Java ArrayList and LinkedList implements the List interface of the Collections framework. However, there exists some difference between them.

LinkedList

Implements List, Queue, and Deque interfaces.

ArrayList

Implements List interface.

Stores 3 values (**previous address**, **data**, and **next address**) in a single position.

Provides the doubly-linked list implementation.

Whenever an element is added, `prev` and `next` address are changed.

To access an element, we need to iterate from the beginning to the element.

Stores a single value in a single position.

Provides a resizable array implementation.

Whenever an element is added, all elements after that position are shifted.

Can randomly access elements using indexes.

Difference between ArrayList and Vector

ArrayList and Vector both implements List interface and maintains insertion order. However, there are many differences between ArrayList and Vector classes that are given below.

ArrayList	Vector
1) ArrayList is not synchronized .	Vector is synchronized .
2) ArrayList increments 50% of current array size if the number of elements exceeds from its capacity.	Vector increments 100% means doubles the array size if the total number of elements exceeds than its capacity.
3) ArrayList is not a legacy class. It is introduced in JDK 1.2.	Vector is a legacy class.
4) ArrayList is fast because it is non-synchronized.	Vector is slow because it is synchronized, i.e., in a multithreading environment, it holds the other threads in runnable or non-runnable state until current thread releases the lock of the object.
5) ArrayList uses the Iterator interface to traverse the elements.	A Vector can use the Iterator interface or Enumeration interface to traverse the elements.

ArrayList vs HashSet.

Sr. No.	Key	ArrayList	HashSet
1	Implementation	ArrayList is the implementation of the list interface.	HashSet on the other hand is the implementation of a set interface.
2	Internal implementation	ArrayList internally implements array for its implementation.	HashSet internally uses Hashmap for its implementation.
3	Order of elements	ArrayList maintains the insertion order i.e order of the object in which they are inserted.	HashSet is an unordered collection and doesn't maintain any order.
4	Duplicates	ArrayList allows duplicate values in its collection.	On other hand duplicate elements are not allowed in Hashset.
5	Index performance	ArrayList uses index for its performance i.e its index based one can retrieve object by calling get(index) or remove objects by calling remove(index)	HashSet is completely based on object also it doesn't provide get() method.
6	Null Allowed	Any number of null value can be inserted in arraylist without any restriction.	On other hand Hashset allows only one null value in its collection, after which no null value is allowed to be added.